

Intro to ML Engineering and Linear Regression: Pre-Reads (Elaborated)

1. Introduction to Machine Learning

What is ML?

Machine Learning is a way to **build systems that learn patterns from data instead of being explicitly programmed with rules.**

Instead of writing:

“If user clicked 3 times → show product A”

You let the system learn:

“Users with similar behavior tend to like product A”

Real Example

- In Netflix:
 - You are not manually assigned shows.
 - ML models analyze your watch history, time spent, genres → predict what you might like.
- In Google:
 - Search ranking is not rule-based.
 - It learns from billions of queries and clicks.

Why ML Matters

- Handles **complex patterns** humans can't manually encode
 - Scales to **millions of users**
 - Continuously improves with more data
-

2. ML vs Classical Programming

Core Difference

Aspect	Classical Programming	Machine Learning
Input	Rules + Data	Data + Labels
Logic	Hand-written	Learned
Output	Deterministic	Probabilistic

Example: Spam Detection

Classical Approach

if "free money" in email:

```
spam = True
```

ML Approach

- Input: Thousands of emails labeled spam/not spam
- Model learns:
 - Word patterns
 - Frequency
 - Context

Key Insight

ML is useful when:

- Rules are **too many** (e.g., fraud detection)
- Rules are **unknown** (e.g., image recognition)

3. SDE Pipeline (Traditional Software Engineering)

Lifecycle

Requirement → Design → Code → Test → Deploy → Maintain

Example: Payment System

- Requirement: Process UPI payments
- Code: Business logic written manually
- Test: Edge cases (failed transactions)
- Deploy: Push to servers

Limitation for ML Problems

- You **can't write rules** for:
 - "What makes a good movie recommendation?"
 - "What is a cat in an image?"
-

4. ML Pipeline (What Changes)

ML pipeline is **data + code + model combined**.

End-to-End Flow

Problem → Data → Features → Model → Evaluation → Deployment → Monitoring

Real Example: House Price Prediction

- Problem: Predict price
- Data: Location, size, bedrooms
- Features: price per sq ft, locality score
- Model: Linear Regression
- Output: Predicted price

Key Difference vs SDE

SDE

ML

Code-centric Data-centric

Static

Continuously
evolving

Debug
code

Debug data +
model

5. Phases of Machine Learning (With Reality)

1. Data Collection

- Sources:
 - Databases
 - APIs
 - Logs
- Example: E-commerce collecting clickstream data

2. Data Preparation (Most Important)

- Missing values
- Outliers
- Encoding categorical variables

Example:

City: Bangalore → [1,0,0]

City: Delhi → [0,1,0]

3. Modeling

- Choose algorithm
- Train model on data

4. Evaluation

- Check performance on unseen data

5. Deployment

- Serve model via API

6. Monitoring

- Detect:
 - Data drift
 - Model degradation
-

6. ML Categorization

ML can be categorized based on:

1. Data Labeling

- Labeled → Supervised
- Unlabeled → Unsupervised

2. Feedback Loop

- With rewards → Reinforcement Learning
-

7. Types of Learning

1. Supervised Learning

- Learn mapping:
 $X \rightarrow y$

Example:

- Input: Resume
 - Output: Selected / Rejected
-

2. Unsupervised Learning

- No labels
- Find hidden structure

Example:

- Customer segmentation:
 - Group users into clusters (high spenders, low spenders)
-

3. Reinforcement Learning

- Agent interacts with environment
- Gets rewards/penalties

Example:

- Self-driving car:
 - Reward: staying in lane
 - Penalty: collision
-

8. ML Task Types (With Examples)

Regression

- Output: Continuous value
 - Example:
 - Predict house price → ₹75 lakhs
-

Classification

- Output: Categories

- Example:
 - Spam vs Not Spam
-

Clustering

- Group similar data
 - Example:
 - Segment customers into 5 groups
-

Recommendation Systems

- Suggest items

Example:

- Amazon:
 - “People who bought this also bought...”
-

Time Series

- Predict future from past

Example:

- Stock prices
 - Demand forecasting
-

Reinforcement Learning Tasks

- Sequential decisions

Example:

- Game AI (e.g., chess engines)
-

9. Linear Regression (Core Concept)

What It Does

Predicts a continuous value using a **linear relationship**.

Example:

- Predict salary from years of experience

Mathematical Representation

$y = w_1x_1 + w_2x_2 + \dots + b$

Meaning of Terms

- xxx: Inputs (features)
- www: Weights (importance)
- bbb: Bias (offset)
- yyy: Prediction


Simple Example



Experience (years)	Salary (₹ LPA)
1	4
2	6
3	8

Model learns:

$$\text{Salary} \approx 2 \times \text{Experience} + 2$$

Geometric Intuition





6

- In 2D → line
- In 3D → plane

- Higher dimensions → hyperplane
-

Why It Matters

- Easy to interpret
 - Fast to train
 - Foundation for:
 - Logistic regression
 - Neural networks (linear layers)
-

10. Problem Statement (Critical Step)

Good Problem Statement Includes:

- Input (X)
- Output (y)
- Constraints
- Success metric

Example

Predict delivery time based on distance, traffic, and weather.

Bad:

“Improve delivery system”

11. Data Notation (Standard Language)

Symb ol	Meaning
------------	---------

X	Input features
---	----------------

y	Target
m	Number of samples
n	Number of features

Example:

- Dataset:
 - 100 houses → $m = 100$
 - 3 features → $n = 3$
-

12. Geometric Intuition (Why It Helps)


Think Visually

- Regression → fitting a line
- Classification → separating data with a boundary
- Clustering → grouping nearby points


Example


- Spam vs Not Spam:
 - Model draws a **decision boundary**


Classification



Regression







6

Final Takeaways

- ML is **data-driven decision making**
- ML pipeline $\neq$ traditional software pipeline
- Linear regression is the **starting point for all ML understanding**
- Real-world ML is:
 - 70% data
 - 20% engineering
 - 10% modeling